

NAME: PASALA NEELIMA

REGNO : 24MDT1064

EXPERIMENT NO : 03, 04

Task : handling missing values

and visualization. And grid &

random search.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

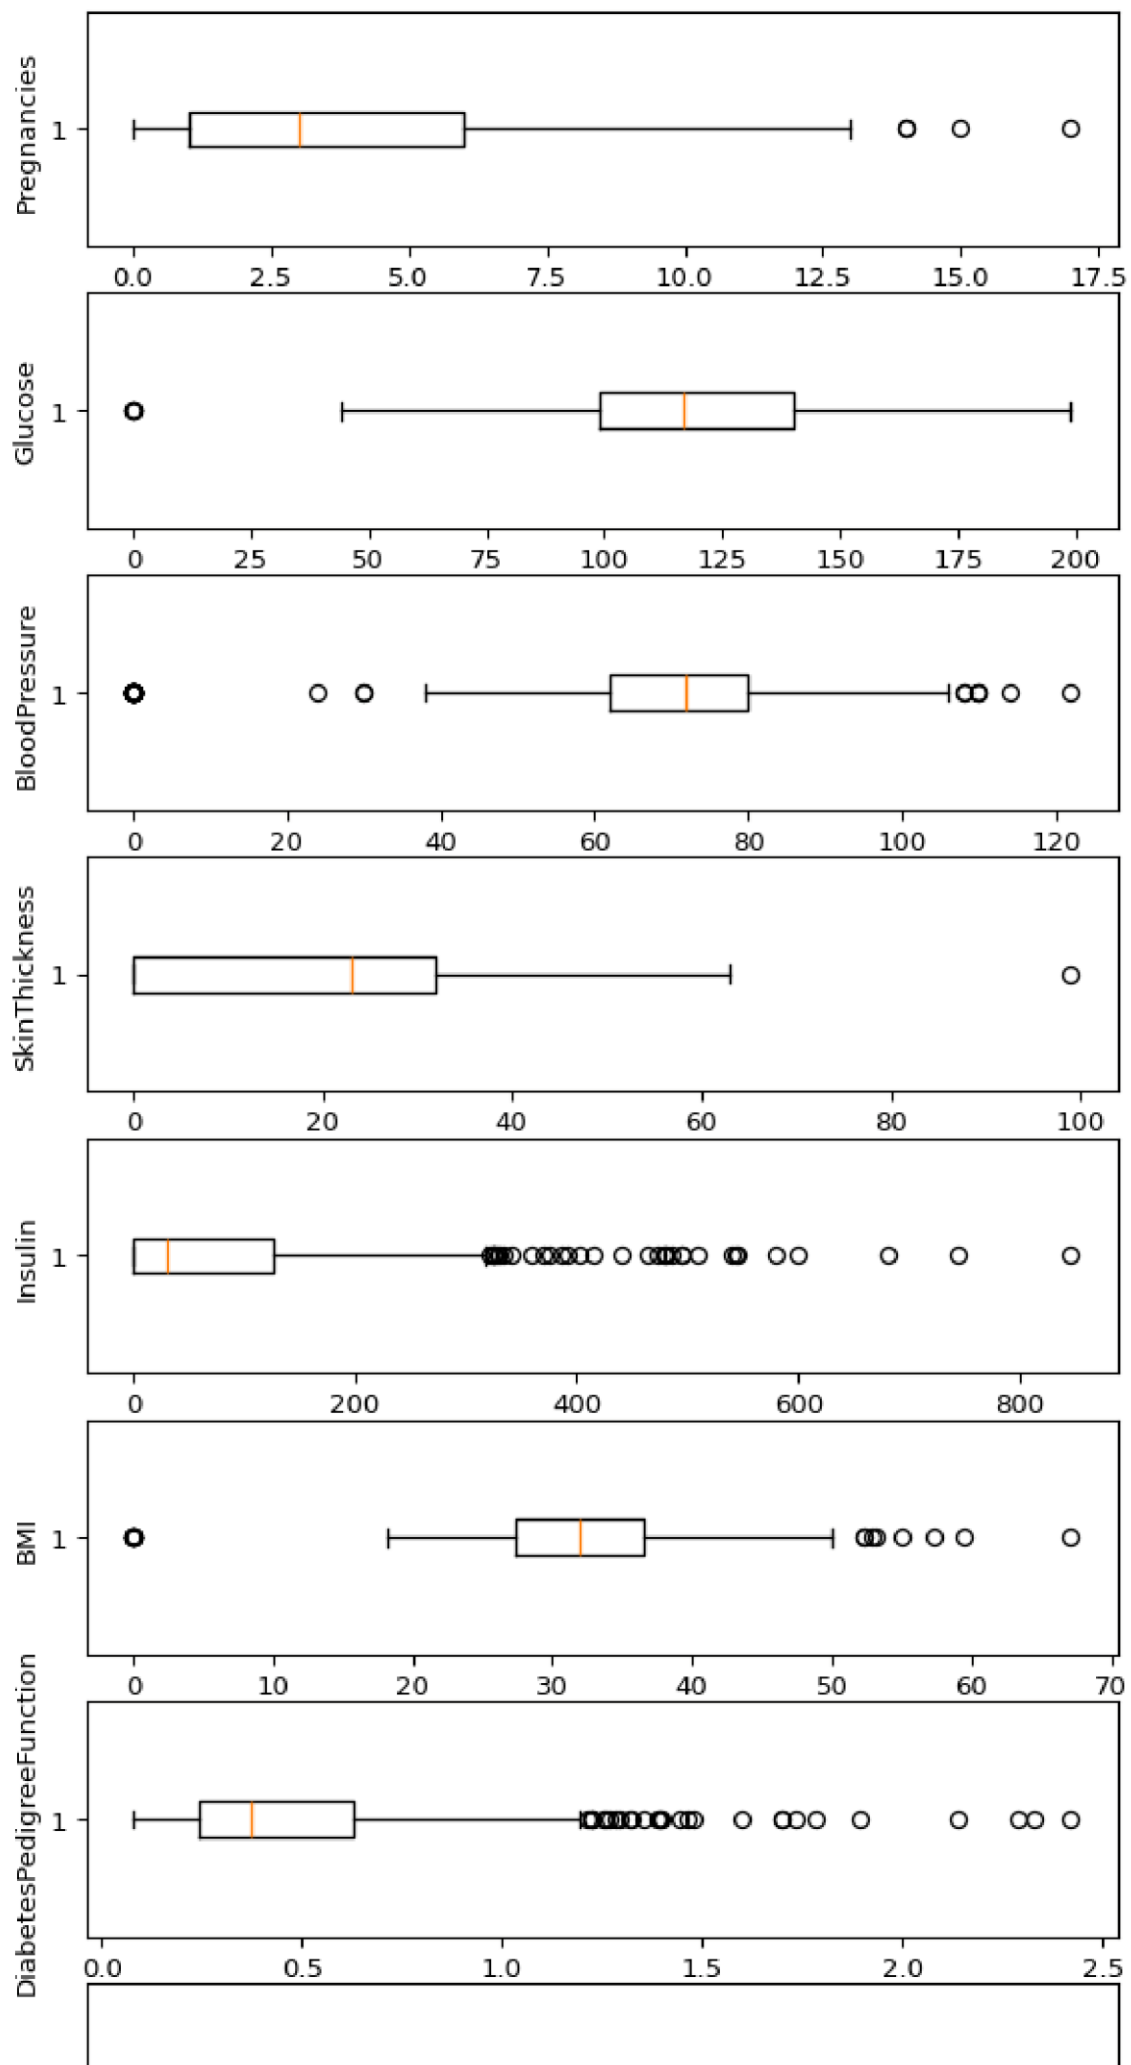
```
df = pd.read_csv("/content/diabetes.csv")
print(df.head())
```


```
➡
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
# BOX PLOT
fig, axes = plt.subplots(9, 1, dpi=95, figsize=(7, 17))
i = 0
# Filter out non-numeric columns before creating boxplots
numeric_cols = df.select_dtypes(include=[np.number]).columns
for col in numeric_cols:
    axes[i].boxplot(df[col], vert=False)
    axes[i].set_ylabel(col)
    i += 1
plt.show()
```






```
#correaltion corr=df.corr()  
sns.heatmap(df.corr(),annot=True,fmt='.2f',cmap="coolwarm"  
) print(corr)
```



	Pregnancies	Glucose	BloodPressure	SkinThickness \
Pregnancies	1.000000	0.129459	0.141282	-0.081672
Glucose	0.129459	1.000000	0.152590	0.057328
BloodPressure	0.141282	0.152590	1.000000	0.207371
SkinThickness	-0.081672	0.057328	0.207371	1.000000
Insulin	-0.073535	0.331357	0.088933	0.436783
BMI	0.017683	0.221071	0.281805	0.392573
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928
Age	0.544341	0.263514	0.239528	-0.113970
Outcome	0.221898	0.466581	0.065068	0.074752

	Insulin	BMI	DiabetesPedigreeFunction \
Pregnancies	-0.073535	0.017683	-0.033523
Glucose	0.331357	0.221071	0.137337
BloodPressure	0.088933	0.281805	0.041265
SkinThickness	0.436783	0.392573	0.183928
Insulin	1.000000	0.197859	0.185071
BMI	0.197859	1.000000	0.140647
DiabetesPedigreeFunction	0.185071	0.140647	1.000000
Age	-0.042163	0.036242	0.033561
Outcome	0.130548	0.292695	0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

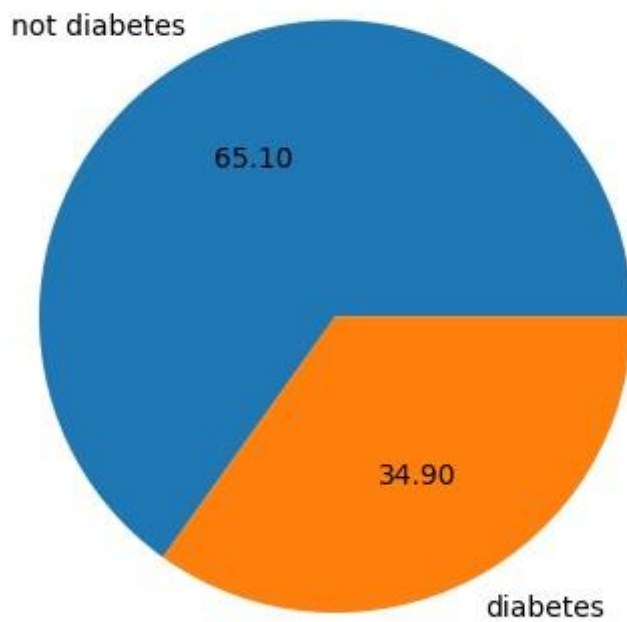


```
plt.pie(df.Outcome.value_counts(), labels=['not
diabetes','diabetes'],autopct='%.2f') plt.title('outcome
proportionality')
```

```
plt.show()
```



outcome proportionality



```
#seperate array into input and output components
```

```
X = df.drop(columns=['Outcome'])
```

```
Y = df.Outcome
```

```
X.head()
```



	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
# drop
```

```
X = df.drop(columns=['Outcome'])
```

```
Y = df.Outcomeprint(X.head())
```



	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age
0	0.627	50

1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler(feature_range=(0, 1))
scaler.fit(X) rescaledX =
scaler.transform(X) print(rescaledX)
```

```
→ [[0.35294118 0.74371859 0.59016393 ... 0.50074516 0.23441503 0.48333333]
    [0.05882353 0.42713568 0.54098361 ... 0.39642325 0.11656704 0.16666667]
    [0.47058824 0.91959799 0.52459016 ... 0.34724292 0.25362938 0.18333333]
    ...
    [0.29411765 0.6080402 0.59016393 ... 0.390462 0.07130658 0.15 ]
    [0.05882353 0.63316583 0.49180328 ... 0.4485842 0.11571307 0.43333333]
    [0.05882353 0.46733668 0.57377049 ... 0.45305514 0.10119556 0.03333333]]
```

```
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
scaler = StandardScaler() rescaledX
= scaler.fit_transform(X)
rescaledX[:5]
```

```
→ array([[ 0.63994726,  0.84832379,  0.14964075,  0.90726993, -0.69289057,
           0.20401277,  0.46849198,  1.4259954 ], [-0.84488505, -
1.12339636, -0.16054575,  0.53090156, -0.69289057, -0.68442195,
-0.36506078, -0.19067191],
 [ 1.23388019,  1.94372388, -0.26394125, -1.28821221, -0.69289057,
-1.10325546,  0.60439732, -0.10558415],
 [-0.84488505, -0.99820778, -0.16054575,  0.15453319,  0.12330164,
-0.49404308, -0.92076261, -1.04154944],
 [-1.14185152,  0.5040552 , -1.50468724,  0.90726993,  0.76583594,
 1.4097456 ,  5.4849091 , -0.0204964 ]])
```

```
import pandas as pd
import numpy as np
```

```
data = {'first test':[100,90,np.nan,95,75,87],
'second test':[30,45,56,np.nan,60,70],
'third test':[np.nan,40,80,98,55,np.nan]}
```

	first test	second test	third test
0	100.0	30.0	NaN
1	90.0	45.0	40.0
2	NaN	56.0	80.0
3	95.0	NaN	98.0

4	75.0	60.0	55.0
5	87.0	70.0	NaN



```
df.dropna()
```

	first test	second test	third test
1	90.0	45.0	40.0
4	75.0	60.0	55.0

```
df = pd.DataFrame(data)
df
```

```
#fillinf missing values with
zero df=df.fillna(0) df
```



	first test	second test	third test
0	100.0	30.0	0.0
1	90.0	45.0	40.0
2	0.0	56.0	80.0
3	95.0	0.0	98.0
4	75.0	60.0	55.0
5	87.0	70.0	0.0

```
df.mean()
```



	0
first test	74.5
second test	43.5
third test	45.5

```
dtype: float64
```

```
df3=df.fillna(df.mean())
print('imputation using mean:')
print(df3)
```



```
imputation using mean:
```

	first test	second test	third test
0	100.0	30.0	0.0
1	90.0	45.0	40.0
2	0.0	56.0	80.0
3	95.0	0.0	98.0
4	75.0	60.0	55.0
5	87.0	70.0	0.0


```
df4=df.fillna(df.median())  
print('imputation using median:')  
print(df4)
```



imputation using median:

	first test	second test	third test
0	100.0	30.0	0.0
1	90.0	45.0	40.0
2	0.0	56.0	80.0
3	95.0	0.0	98.0
4	75.0	60.0	55.0
5	87.0	70.0	0.0

```
df5=df.ffill()  
print('imputationusing next  
value:') print(df5)
```

NAME: PASALA NEELIMA

REGNO : 24MDT1064 SLOT:

LU5,U6

EXPERIMENT NO : 04 DATE

04-02-2025

```
import pandas as pd
df = pd.read_csv("/content/iris.csv")
print(df.head())
```

```

sepal_length  sepal_width  petal_length  petal_width  species
0      5.1         3.5         1.4         0.2    setosa
1      4.9         3.0         1.4         0.2    setosa
2      4.7         3.2         1.3         0.2    setosa
3      4.6         3.1         1.5         0.2    setosa
4      5.0         3.6         1.4         0.2    setosa

```

```
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score
```

```
iris = datasets.load_iris()
x = iris.data
y = iris.target
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
param_grid = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': ['scale', 'auto']
}
```

```
# setup the gridsearch cv
```

```
grid_search = GridSearchCV(estimator=SVC, param_grid=param_grid, cv=5, verbose=3, n_jobs=-1, scoring='accuracy')
grid_search.fit(x_train, y_train)
```

```
Fitting 5 folds for each of 24 candidates, totalling 120 fits
```

```

GridSearchCV
└─ best_estimator_:
    SVC
    └─ SVC

```

```
# GET THE BEST PARAMETERS AND BESTSCORE
```

```
best_params = grid_search.best_params_
best_score = grid_search.best_score_
print("best parameters:", best_params)
print("bestcross-validation score:", best_score)
```

```
best parameters: {'C': 1, 'gamma': 'scale', 'kernel': 'linear'} bestcross-validation score: 0.9583333333333334
```

```
best_model = grid_search.best_estimator_
y_pred = best_model.predict(x_test)
```

```
#print classification report
print("classification report:")
print(classification_report(y_test, y_pred))
```

```

classification report:
              precision
recall  f1-score  support
0      1.00      1.00      10
1      1.00      1.00       9

```

2	1.00	1.00	1.00	11	
accuracy				1.00	30
macro avg	1.00	1.00	1.00	30	
weighted avg	1.00	1.00	1.00	30	

```
accuracy = accuracy_score(y_test,y_pred)
```

```
print("accuracy on test set:",accuracy)
```

```
accuracy on test set: 1.0
```

```
import numpy as np from sklearn import datasets from
sklearn.model_selection import train_test_split,RandomizedSearchCV from
sklearn.svm import SVC
from sklearn.metrics import classification_report,accuracy_score
from scipy.stats import uniform, randint
```

```
iris = datasets.load_iris()
x = iris.data y =
iris.target
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
model = SVC()
```

```
param_distributions = {
    'C' : uniform(0.1,100),
    'kernel':['linear','rbf','poly'],
    'gamma' :['scale','auto'] + list(uniform(0.001,1).rvs(size=10))
}
```

```
} Start coding or generate with AI.
```


```
# setup the randomizedsearchcv
random_search = RandomizedSearchCV(estimator=model, param_distributions=param_distributions, n_iter=100, cv=5, verbose=2, n_jobs=-1, scor
random_search.fit(x_train, y_train)
```

```
# GET THE BEST PARAMETERS AND BESTSCORE
best_params_random = random_search.best_params_
best_score_random = random_search.best_score_
print("best parameters:", best_params_random)
print("bestcross-validation score:",best_score_random)
best_model_random = random_search.best_estimator_
y_pred_random = best_model_random.predict(x_test)
#print classification report print("classification
report:")
print(classification_report(y_test,y_pred_random))
accuracy_random = accuracy_score(y_test,y_pred_random)
print("accuracy on test set:",accuracy_random)
```

```
Fitting 5 folds for each of 100 candidates, totalling 500 fits
best parameters: {'C': 19.69828624191452, 'gamma': 0.07132602069540483, 'kernel':
'rbf'} bestcross-validation score: 0.9666666666666668 classification report:
precision recall f1-score support
```

0	1.00	1.00	1.00	10				
1	1.00	1.00	1.00	9	2	1.00	1.00	1.00
11								
accuracy				1.00	30			
macro avg	1.00	1.00	1.00	30				
weighted avg	1.00	1.00	1.00	30				
accuracy on test set: 1.0								

```
random_search.fit(x_train, y_train)
```



```
Fitting 5 folds for each of 100 candidates, totalling 500 fits
```